

```

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('whitegrid')
plt.style.use("fivethirtyeight")
%matplotlib inline

# Para poder leer y trabajar con los precios de las acciones
from pandas_datareader.data import DataReader
import yfinance as yf
from pandas_datareader import data as pdr

yf.pdr_override()

# Para la fecha de los datos
from datetime import datetime

# Lista con el nombre de las acciones que vamos a utilizar
tech_list = ['AAPL', 'GOOG', 'MSFT', 'META']

# Definimos qué fechas queremos de los datos
end = datetime.now()
start = datetime(end.year - 1, end.month, end.day)

for stock in tech_list:
    globals()[stock] = yf.download(stock, start = "2012-01-01", end =
"2024-01-05")

company_list = [AAPL, GOOG, MSFT, META]
company_name = ["APPLE", "GOOGLE", "MICROSOFT", "META"]

for company, com_name in zip(company_list, company_name):
    company["company_name"] = com_name

df = pd.concat(company_list, axis=0)
df.tail(10)

```

```

[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed

```

	Open	High	Low	Close	Adj Close
\					
Date					
2023-12-20	348.649994	354.959991	347.790009	349.279999	348.909790

2023-12-21	352.980011	356.410004	349.209991	354.089996	353.714691
2023-12-22	355.579987	357.200012	351.220001	353.390015	353.015472
2023-12-26	354.989990	356.980011	353.450012	354.829987	354.453918
2023-12-27	356.070007	359.000000	355.309998	357.829987	357.450714
2023-12-28	359.700012	361.899994	357.809998	358.320007	357.940216
2023-12-29	358.989990	360.000000	351.820007	353.959991	353.584839
2024-01-02	351.320007	353.160004	340.010010	346.290009	345.922974
2024-01-03	344.980011	347.950012	343.179993	344.470001	344.104889
2024-01-04	344.500000	348.149994	343.399994	347.119995	346.752075

	Volume	company_name
Date		
2023-12-20	16369900	META
2023-12-21	15289600	META
2023-12-22	11764200	META
2023-12-26	9898600	META
2023-12-27	13207900	META
2023-12-28	11798800	META
2023-12-29	14980500	META
2024-01-02	19042200	META
2024-01-03	15451100	META
2024-01-04	12099900	META


#Graficamos los precios de las acciones de este último año

```
plt.figure(figsize=(15, 10))
plt.subplots_adjust(top=1.25, bottom=1.2)

for i, company in enumerate(company_list, 1):
    plt.subplot(2, 2, i)
    company['Adj Close'].plot(linewidth = 1)

    plt.xlabel(None)
    plt.title(f"Cierre ajustado de {tech_list[i - 1]}")

plt.tight_layout()
```



```
import yfinance as yf
import matplotlib.pyplot as plt

# Lista de tickers de las empresas
tickers = ["AAPL", "GOOG", "MSFT", "META"]

# Obtener los datos de precios de las acciones para el último año
start_date = "2012-01-01"
end_date = "2024-05-01"

data = yf.download(tickers, start=start_date, end=end_date)

# Calcular los retornos diarios
returns = data["Adj Close"].pct_change()

# Calcular la volatilidad para cada diez días
volatility_10_days = returns.rolling(window=10).std() * (252 ** 0.5)
# Ventana de 10 días, asumiendo 252 días hábiles en un año

# Crear subgráficos
fig, axes = plt.subplots(2, 2, figsize=(14, 12))

# Graficar la volatilidad para cada empresa en subgráficos separados
for i, ticker in enumerate(tickers):
    row = i // 2
```

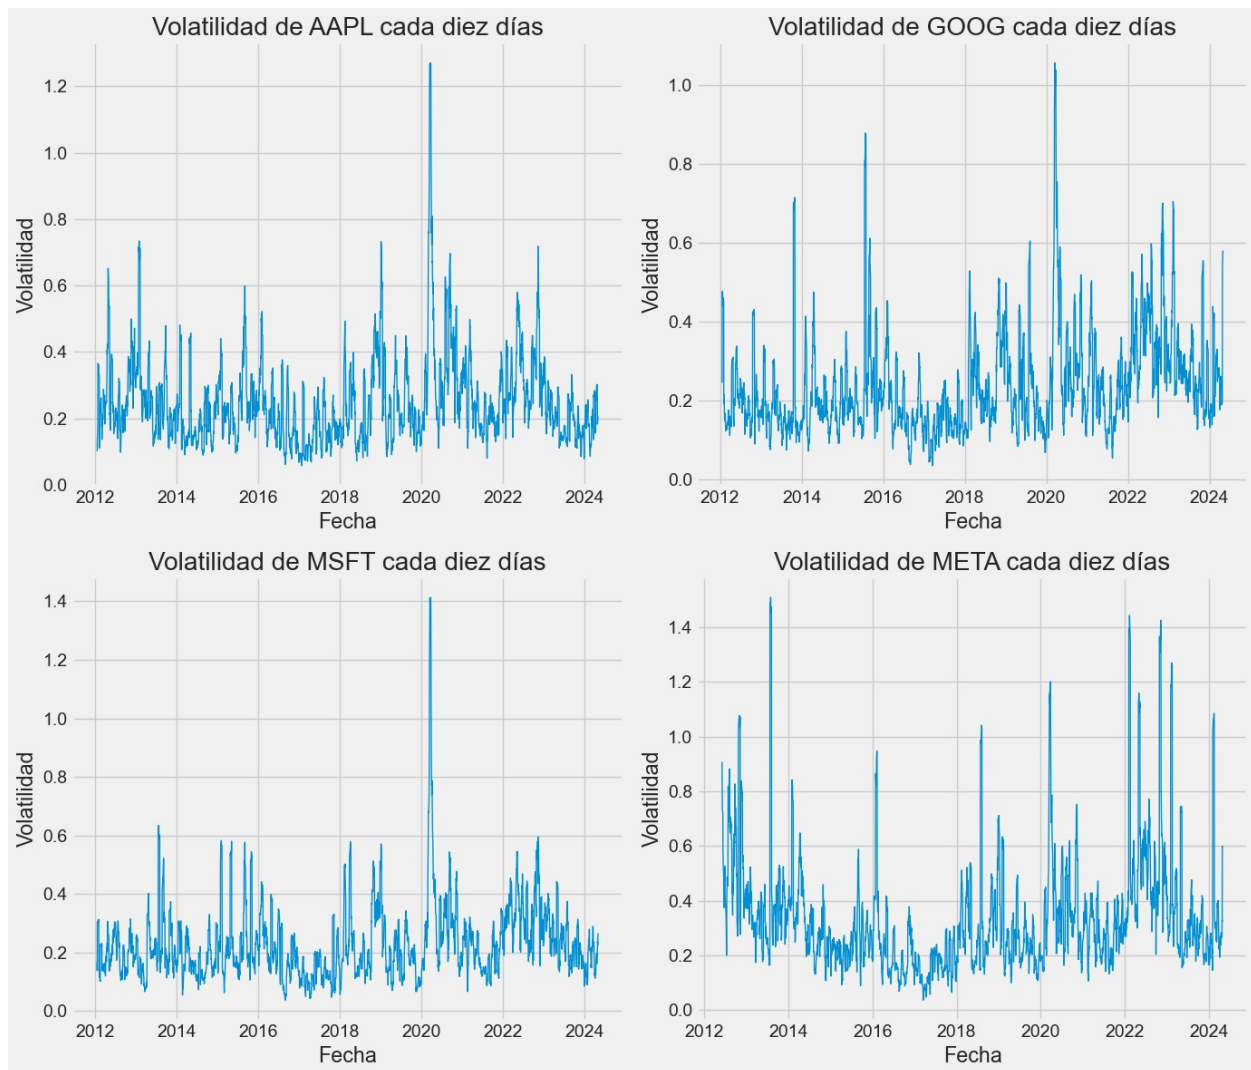
```

col = i % 2
axes[row, col].plot(volatility_10_days[ticker], linewidth=1)
axes[row, col].set_title(f'Volatilidad de {ticker} cada diez
días')
axes[row, col].set_xlabel('Fecha')
axes[row, col].set_ylabel('Volatilidad')

plt.tight_layout()
plt.show()

[*****100%*****] 4 of 4 completed

```



```

import yfinance as yf

# Lista de tickers de las empresas
tickers = ["AAPL", "GOOG", "MSFT", "META"]

```

```

# Obtener los datos de precios de las acciones para el último año
start_date = "2023-05-01"
end_date = "2024-05-01"

data = yf.download(tickers, start=start_date, end=end_date)

# Calcular los retornos diarios
returns = data["Adj Close"].pct_change()

# Calcular la volatilidad anualizada asumiendo 252 días hábiles
volatility = returns.std() * (252 ** 0.5)

print("Volatilidad anualizada:")
for ticker, vol in zip(tickers, volatility):
    print(f"{ticker}: {vol:.2f}")

import matplotlib.pyplot as plt


# Lista de nombres de empresas
companies = ["Apple (AAPL)", "Alphabet (GOOG)", "Microsoft (MSFT)",
            "Meta Platforms (META)"]

# Crear un gráfico de barras para visualizar la volatilidad
plt.figure(figsize=(9, 6))
plt.bar(companies, volatility, color='blue', width=0.45)
plt.title('Volatilidad de las acciones en el último año')
plt.xlabel('Empresa')
plt.ylabel('Volatilidad anualizada')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

[*****100%*****] 4 of 4 completed

Volatilidad anualizada:
AAPL: 0.20
GOOG: 0.29
MSFT: 0.36
META: 0.21

```



```
ma_day = [10, 20, 50]

for ma in ma_day:
    for company in company_list:
        column_name = f"MA for {ma} days"
        company[column_name] = company['Adj Close'].rolling(ma).mean()

fig, axes = plt.subplots(nrows=2, ncols=2)
fig.set_figheight(10)
fig.set_figwidth(15)

AAPL[['Adj Close', 'MA for 10 days', 'MA for 20 days', 'MA for 50
days']].plot(ax=axes[0,0])
axes[0,0].set_title('APPLE')

GOOG[['Adj Close', 'MA for 10 days', 'MA for 20 days', 'MA for 50
days']].plot(ax=axes[0,1])
axes[0,1].set_title('GOOGLE')

MSFT[['Adj Close', 'MA for 10 days', 'MA for 20 days', 'MA for 50
days']].plot(ax=axes[1,0])
axes[1,0].set_title('MICROSOFT')


META[['Adj Close', 'MA for 10 days', 'MA for 20 days', 'MA for 50
```

```

days']].plot(ax=axes[1,1])
axes[1,1].set_title('META')

fig.tight_layout()

```



```

# Usamos pct_change para encontrar el porcentaje de cambio en el
# precio cada día
for company in company_list:
    company['Daily Return'] = company['Adj Close'].pct_change()

# Graficamos
fig, axes = plt.subplots(nrows=2, ncols=2)
fig.set_figheight(10)
fig.set_figwidth(15)

AAPL['Daily Return'].plot(ax=axes[0,0], legend=True, linestyle='--',
marker='.', linewidth=1)
axes[0,0].set_title('APPLE')

GOOG['Daily Return'].plot(ax=axes[0,1], legend=True, linestyle='--',
marker='.', linewidth=1)
axes[0,1].set_title('GOOGLE')

MSFT['Daily Return'].plot(ax=axes[1,0], legend=True, linestyle='--',

```

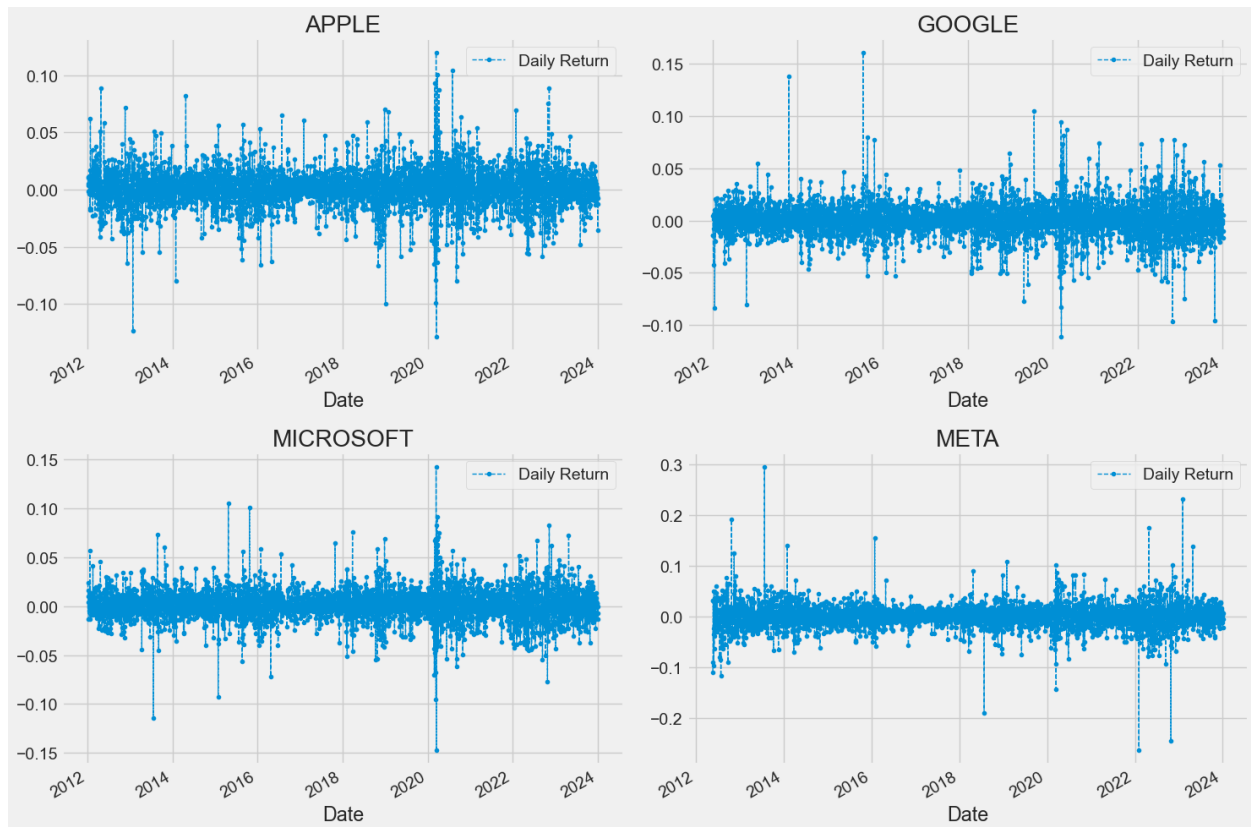
```

marker='.', linewidth=1)
axes[1,0].set_title('MICROSOFT')

META['Daily Return'].plot(ax=axes[1,1], legend=True, linestyle='--',
marker='.', linewidth=1)
axes[1,1].set_title('META')

fig.tight_layout()

```



```


rets = tech_rets.dropna()

area = np.pi * 20

plt.figure(figsize=(10, 8))
plt.scatter(rets.mean(), rets.std(), s=area)
plt.xlabel('Retorno esperado')
plt.ylabel('Riesgo')

for label, x, y in zip(rets.columns, rets.mean(), rets.std()):
    plt.annotate(label, xy=(x, y), xytext=(50, 50), textcoords='offset
points', ha='right', va='bottom',
                arrowprops=dict(arrowstyle='--', color='blue',
connectionstyle='arc3,rad=-0.3'))

```

Predicción con LSTM

```
# Cargamos los datos de las acciones que queremos analizar y predecir
df = pdr.get_data_yahoo('AAPL', start='2012-01-01',
end=datetime.now())
```

df

```
[*****100%*****] 1 of 1 completed
```

	Open	High	Low	Close	Adj Close
\ Date					
2012-01-03	14.621429	14.732143	14.607143	14.686786	12.416983
2012-01-04	14.642857	14.810000	14.617143	14.765714	12.483713
2012-01-05	14.819643	14.948214	14.738214	14.929643	12.622309
2012-01-06	14.991786	15.098214	14.972143	15.085714	12.754258

2012-01-09	15.196429	15.276786	15.048214	15.061786	12.734027
...
2024-05-13	185.440002	187.100006	184.619995	186.279999	186.279999
2024-05-14	187.509995	188.300003	186.289993	187.429993	187.429993
2024-05-15	187.910004	190.649994	187.369995	189.720001	189.720001
2024-05-16	190.470001	191.100006	189.660004	189.839996	189.839996
2024-05-17	189.380005	190.809998	189.220001	189.869995	189.869995

	Volume
Date	
2012-01-03	302220800
2012-01-04	260022000
2012-01-05	271269600
2012-01-06	318292800
2012-01-09	394024400
...	...
2024-05-13	72044800
2024-05-14	52393600
2024-05-15	70400000
2024-05-16	52845200
2024-05-17	37956559

[3114 rows x 6 columns]

```
#Hacemos una gráfica con el histórico de precios
plt.figure(figsize=(16,6))
plt.title('Histórico de precios AAPL')
plt.plot(df['Adj Close'], linewidth = 1)
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre ajustado USD ($)', fontsize=18)
plt.show()
```



```
# Creamos un nuevo DataFrame con únicamente la columna Adj Close
data = df.filter(['Adj Close'])
# Convertimos el dataframe en un numpy array
dataset = data.values
# Número de filas para entrenar el modelo
training_data_len = int(np.ceil( len(dataset) * .95 ))

training_data_len

2959

# Escalamos y normalizamos los datos
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler(feature_range=(0,1))
scaled_data = scaler.fit_transform(dataset)

scaled_data

array([[0.00242952],
       [0.00278902],
       [0.00353566],
       ...,
       [0.95760493],
       [0.95825137],
       [0.95841298]])

# Creamos el data set para entrenamiento con los datos escalados
train_data = scaled_data[0:int(training_data_len), :]
# Separamos los datos entre x_train y y_train
x_train = []
y_train = []

for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:i, 0])
```

```

y_train.append(train_data[i, 0])
if i<= 61:
    print(x_train)
    print(y_train)
    print()

# Convertimos x_train y y_train en numpy arrays
x_train, y_train = np.array(x_train), np.array(y_train)

# Cambiamos la forma de los datos
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))

[array([0.00242952, 0.00278902, 0.00353566, 0.00424651, 0.00413752,
        0.00438316, 0.00427091, 0.0040822 , 0.0038252 , 0.00462065,
        0.00533801, 0.00511677, 0.00390492, 0.00506149, 0.00392281,
        0.00819283, 0.00786259, 0.00829366, 0.00922574, 0.0097902 ,
        0.00974301, 0.00956897, 0.0103107 , 0.01100855, 0.01179914,
        0.01307605, 0.01575844, 0.01579909, 0.01729238, 0.01840826,
        0.01649044, 0.01722893, 0.01721427, 0.01928502, 0.01899058,
        0.01953552, 0.02051477, 0.02105974, 0.02262947, 0.023773 ,
        0.02410324, 0.02421869, 0.02226349, 0.02179174, 0.02186169,
        0.02369983, 0.02421709, 0.02532813, 0.02794702, 0.03144112,
        0.03078715, 0.03078883, 0.033315 , 0.03410559, 0.03354275,
        0.03302875, 0.03249354, 0.0342715 , 0.03549148, 0.0360023 ])]
[0.034739983147894324]

[array([0.00242952, 0.00278902, 0.00353566, 0.00424651, 0.00413752,
        0.00438316, 0.00427091, 0.0040822 , 0.0038252 , 0.00462065,
        0.00533801, 0.00511677, 0.00390492, 0.00506149, 0.00392281,
        0.00819283, 0.00786259, 0.00829366, 0.00922574, 0.0097902 ,
        0.00974301, 0.00956897, 0.0103107 , 0.01100855, 0.01179914,
        0.01307605, 0.01575844, 0.01579909, 0.01729238, 0.01840826,
        0.01649044, 0.01722893, 0.01721427, 0.01928502, 0.01899058,
        0.01953552, 0.02051477, 0.02105974, 0.02262947, 0.023773 ,
        0.02410324, 0.02421869, 0.02226349, 0.02179174, 0.02186169,
        0.02369983, 0.02421709, 0.02532813, 0.02794702, 0.03144112,
        0.03078715, 0.03078883, 0.033315 , 0.03410559, 0.03354275,
        0.03302875, 0.03249354, 0.0342715 , 0.03549148, 0.0360023 ]),
array([0.00278902, 0.00353566, 0.00424651, 0.00413752, 0.00438316,
        0.00427091, 0.0040822 , 0.0038252 , 0.00462065, 0.00533801,
        0.00511677, 0.00390492, 0.00506149, 0.00392281, 0.00819283,
        0.00786259, 0.00829366, 0.00922574, 0.0097902 , 0.00974301,
        0.00956897, 0.0103107 , 0.01100855, 0.01179914, 0.01307605,
        0.01575844, 0.01579909, 0.01729238, 0.01840826, 0.01649044,
        0.01722893, 0.01721427, 0.01928502, 0.01899058, 0.01953552,
        0.02051477, 0.02105974, 0.02262947, 0.023773 , 0.02410324,
        0.02421869, 0.02226349, 0.02179174, 0.02186169, 0.02369983,
        0.02421709, 0.02532813, 0.02794702, 0.03144112, 0.03078715,
        0.03078883, 0.033315 , 0.03410559, 0.03354275, 0.03302875,
        0.03249354, 0.0342715 , 0.03549148, 0.0360023 , 0.03473998)]]
```

```
[0.034739983147894324, 0.03306286956459417]
```

```
from keras.models import Sequential
from keras.layers import LSTM, Dense

# Creamos el modelo LSTM
model = Sequential()

# Añadimos capas LSTM con 50 unidades
model.add(LSTM(50, return_sequences=True,
input_shape=(x_train.shape[1], 1)))
model.add(LSTM(50, return_sequences=False))
# Añadimos una capa densa con 25 unidades
model.add(Dense(25))
# Añadir una capa de salida densa con 1 unidad
model.add(Dense(1))

# Compilamos el modelo
model.compile(optimizer='adam', loss='mean_squared_error')

# Entrenamos el modelo
model.fit(x_train, y_train, batch_size=1, epochs=1)

# Creamos el conjunto de datos de prueba
test_data = scaled_data[training_data_len - 60: , :]
# Creamos los conjuntos de datos x_test y y_test
x_test = []
y_test = dataset[training_data_len:, :]
for i in range(60, len(test_data)):
    x_test.append(test_data[i-60:i, 0])

# Convertimos a un numpy array
x_test = np.array(x_test)

# Cambiamos la forma de los datos
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))

# Obtención de las predicciones
predictions = model.predict(x_test)
predictions = scaler.inverse_transform(predictions)

# Calculamos el RMSE
rmse=np.sqrt(np.mean(((predictions- y_test)**2)))
print("Root Mean Square Error:",rmse)

# Calcular el NRMSE
range_y = np.max(y_test) - np.min(y_test)
nrmse = rmse / range_y
print("Normalized Root Mean Square Error (NRMSE):", nrmse)
```

```

# Calculamos el MAE
mae = np.mean(np.abs(predictions - y_test))
print("Mean Absolute Error:", mae)

# Calculamos el coeficiente de determinación (R cuadrado)
from sklearn.metrics import r2_score

r2 = r2_score(y_test, predictions)
print("Coeficiente de determinación (R cuadrado):", r2)

c:\Users\Usuario\anaconda3\Lib\site-packages\keras\src\layers\rnn\
rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

2899/2899 ————— 25s 8ms/step - loss: 0.0021
5/5 ————— 0s 57ms/step
Root Mean Square Error: 4.225860178279189
Normalized Root Mean Square Error (NRMSE): 0.1287860834806857
Mean Absolute Error: 3.4758476995652723
Coeficiente de determinación (R cuadrado): 0.7764756512759355


# Graficamos los datos
train = data[:training_data_len]
valid = data[training_data_len:]
valid['Predictions'] = predictions

plt.figure(figsize=(16,6))
plt.title('Apple')
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre USD ($)', fontsize=18)
plt.plot(train['Adj Close'], linewidth = 1)
plt.plot(valid[['Adj Close', 'Predictions']], linewidth = 1)
plt.legend(['Datos de entrenamiento', 'Datos de prueba',
'Predicciones'], loc='lower right')
plt.show()

C:\Users\Usuario\AppData\Local\Temp\ipykernel_14880\3264202307.py:4:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#
returning-a-view-versus-a-copy
  valid['Predictions'] = predictions

```




```
# Separar los datos de prueba
valid = data[training_data_len:]
valid['Predictions'] = predictions

# Graficar los datos de prueba y las predicciones
plt.figure(figsize=(16, 6))
plt.title('Apple')
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre USD ($)', fontsize=18)
plt.plot(valid['Adj Close'], linewidth=1, label='Datos de prueba')
plt.plot(valid['Predictions'], linewidth=1, label='Predicciones')
plt.legend(loc='lower right')
plt.show()
```

C:\Users\Usuario\AppData\Local\Temp\ipykernel_14880\682196135.py:3:
 SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 valid['Predictions'] = predictions



```
# Cargamos los datos de las acciones que queremos analizar y predecir
df = pdr.get_data_yahoo('GOOG', start='2012-01-01',
end=datetime.now())
```

df

[*****100%*****] 1 of 1 completed

\	Open	High	Low	Close	Adj Close
Date					
2012-01-03	16.262545	16.641375	16.248346	16.573130	16.573130
2012-01-04	16.563665	16.693678	16.453827	16.644611	16.644611
2012-01-05	16.491436	16.537264	16.344486	16.413727	16.413727
2012-01-06	16.417213	16.438385	16.184088	16.189817	16.189817
2012-01-09	16.102144	16.114599	15.472754	15.503389	15.503389
...
2024-05-13	165.847000	170.949997	165.759995	170.899994	170.899994
2024-05-14	171.589996	172.779999	170.419998	171.929993	171.929993
2024-05-15	172.300003	174.046005	172.029999	173.880005	173.880005
2024-05-16	174.600006	176.339996	174.050003	175.429993	175.429993
2024-05-17	175.647003	177.490005	174.979996	177.289993	177.289993

Volume


```

Date
2012-01-03    147611217
2012-01-04    114989399
2012-01-05    131808205
2012-01-06    108119746
2012-01-09    233776981
...
2024-05-13    19648600
2024-05-14    18729500
2024-05-15    20958200
2024-05-16    17247300
2024-05-17    16071736

```

```
[3114 rows x 6 columns]
```

```

# Hacemos una gráfica con el histórico de precios
plt.figure(figsize=(16,6))
plt.title('Histórico de precios GOOG')
plt.plot(df['Adj Close'], linewidth = 1)
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre ajustado USD ($)', fontsize=18)
plt.show()

```



```

# Creamos un nuevo DataFrame con únicamente la columna Adj Close
data = df.filter(['Adj Close'])
# Convertimos el dataframe en un numpy array
dataset = data.values
# Número de filas para entrenar el modelo
training_data_len = int(np.ceil( len(dataset) * .95 ))

training_data_len

2959

```

```

# Escalamos y normalizamos los datos
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler(feature_range=(0,1))
scaled_data = scaler.fit_transform(dataset)

scaled_data

array([[0.01621556],
       [0.01665312],
       [0.01523982],
       ...,
       [0.97912669],
       [0.98861451],
       [1.         ]])

# Creamos el data set para entrenamiento con los datos escalados
train_data = scaled_data[0:int(training_data_len), :]
# Separamos los datos entre x_train y y_train
x_train = []
y_train = []

for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:i, 0])
    y_train.append(train_data[i, 0])
    if i<= 61:
        print(x_train)
        print(y_train)
        print()

# Convertimos x_train y y_train en numpy arrays
x_train, y_train = np.array(x_train), np.array(y_train)

# Cambiamos la forma de los datos
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))

[array([0.01621556, 0.01665312, 0.01523982, 0.01386922, 0.00966744,
        0.00977111, 0.01020105, 0.01076209, 0.01005316, 0.01060049,
        0.01126064, 0.01227602, 0.00410725, 0.0040356 , 0.00333581,
        0.00159168, 0.00137975, 0.00319097, 0.00284184, 0.00321079,
        0.00332056, 0.00397309, 0.00568368, 0.00762906, 0.00727536,
        0.00774493, 0.00799039, 0.00714424, 0.00810321, 0.00773121,
        0.00709088, 0.00723724, 0.00695062, 0.00837763, 0.00745373,
        0.00717473, 0.00775255, 0.0076626 , 0.00904693, 0.00902559,
        0.00965829, 0.00948296, 0.00841575, 0.0069994 , 0.00727993,
        0.00733177, 0.00628132, 0.00702837, 0.00895393, 0.00868103,
        0.00946468, 0.01006079, 0.01142377, 0.01134907, 0.01233852,
        0.01326395, 0.01273645, 0.01376402, 0.01341184, 0.01474433])]
[0.013623760824461206]

```

```
[array([0.01621556, 0.01665312, 0.01523982, 0.01386922, 0.00966744,
        0.00977111, 0.01020105, 0.01076209, 0.01005316, 0.01060049,
        0.01126064, 0.01227602, 0.00410725, 0.0040356 , 0.00333581,
        0.00159168, 0.00137975, 0.00319097, 0.00284184, 0.00321079,
        0.00332056, 0.00397309, 0.00568368, 0.00762906, 0.00727536,
        0.00774493, 0.00799039, 0.00714424, 0.00810321, 0.00773121,
        0.00709088, 0.00723724, 0.00695062, 0.00837763, 0.00745373,
        0.00717473, 0.00775255, 0.0076626 , 0.00904693, 0.00902559,
        0.00965829, 0.00948296, 0.00841575, 0.0069994 , 0.00727993,
        0.00733177, 0.00628132, 0.00702837, 0.00895393, 0.00868103,
        0.00946468, 0.01006079, 0.01142377, 0.01134907, 0.01233852,
        0.01326395, 0.01273645, 0.01376402, 0.01341184, 0.01474433]),
array([0.01665312, 0.01523982, 0.01386922, 0.00966744, 0.00977111,
        0.01020105, 0.01076209, 0.01005316, 0.01060049, 0.01126064,
        0.01227602, 0.00410725, 0.0040356 , 0.00333581, 0.00159168,
        0.00137975, 0.00319097, 0.00284184, 0.00321079, 0.00332056,
        0.00397309, 0.00568368, 0.00762906, 0.00727536, 0.00774493,
        0.00799039, 0.00714424, 0.00810321, 0.00773121, 0.00709088,
        0.00723724, 0.00695062, 0.00837763, 0.00745373, 0.00717473,
        0.00775255, 0.0076626 , 0.00904693, 0.00902559, 0.00965829,
        0.00948296, 0.00841575, 0.0069994 , 0.00727993, 0.00733177,
        0.00628132, 0.00702837, 0.00895393, 0.00868103, 0.00946468,
        0.01006079, 0.01142377, 0.01134907, 0.01233852, 0.01326395,
        0.01273645, 0.01376402, 0.01341184, 0.01474433, 0.01362376)])
[0.013623760824461206, 0.012530624495400211]
```

```
from keras.models import Sequential
from keras.layers import LSTM, Dense

# Creamos el modelo LSTM
model = Sequential()

# Añadimos capas LSTM con 50 unidades
model.add(LSTM(50, return_sequences=True,
input_shape=(x_train.shape[1], 1)))
model.add(LSTM(50, return_sequences=False))
# Añadimos una capa densa con 25 unidades
model.add(Dense(25))
# Añadir una capa de salida densa con 1 unidad
model.add(Dense(1))

# Compilamos el modelo
model.compile(optimizer='adam', loss='mean_squared_error')

# Entrenamos el modelo
model.fit(x_train, y_train, batch_size=1, epochs=1)

# Creamos el conjunto de datos de prueba
```

```

test_data = scaled_data[training_data_len - 60: , :]
# Creamos los conjuntos de datos x_test y y_test
x_test = []
y_test = dataset[training_data_len:, :]
for i in range(60, len(test_data)):
    x_test.append(test_data[i-60:i, 0])

# Convertimos a un numpy array
x_test = np.array(x_test)

# Cambiamos la forma de los datos
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))

# Obtención de las predicciones
predictions = model.predict(x_test)
predictions = scaler.inverse_transform(predictions)

# Calculamos el RMSE
rmse=np.sqrt(np.mean(((predictions- y_test)**2)))
print("Root Mean Square Error:",rmse)

# Calcular el NRMSE
range_y = np.max(y_test) - np.min(y_test)
nrmse = rmse / range_y
print("Normalized Root Mean Square Error (NRMSE):", nrmse)

# Calculamos el MAE
mae = np.mean(np.abs(predictions - y_test))
print("Mean Absolute Error:", mae)

# Calculamos el coeficiente de determinación (R cuadrado)
from sklearn.metrics import r2_score

r2 = r2_score(y_test, predictions)
print("Coeficiente de determinación (R cuadrado):", r2)

```

```

c:\Users\Usuario\anaconda3\Lib\site-packages\keras\src\layers\rnn\
rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

```

```

2899/2899 ————— 25s 8ms/step - loss: 0.0018
WARNING:tensorflow:5 out of the last 11 calls to <function
TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_dist
ributed at 0x000001F46CF6D4E0> triggered tf.function retracing.
Tracing is expensive and the excessive number of tracings could be due
to (1) creating @tf.function repeatedly in a loop, (2) passing tensors
with different shapes, (3) passing Python objects instead of tensors.
For (1), please define your @tf.function outside of the loop. For (2),

```

@tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

1/5 ————— 0s 212ms/step WARNING:tensorflow:5 out of the last 11 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x000001F46CF6D4E0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

5/5 ————— 0s 57ms/step
 Root Mean Square Error: 3.914599464549233
 Normalized Root Mean Square Error (NRMSE): 0.07264056528278061
 Mean Absolute Error: 3.022977472120716
 Coeficiente de determinación (R cuadrado): 0.892356244385851

Graficamos los datos

```
train = data[:training_data_len]
valid = data[training_data_len:]
valid['Predictions'] = predictions
```

```
plt.figure(figsize=(16,6))
plt.title('Google')
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre USD ($)', fontsize=18)
plt.plot(train['Adj Close'], linewidth = 1)
plt.plot(valid[['Adj Close', 'Predictions']], linewidth = 1)
plt.legend(['Datos de entrenamiento', 'Datos de prueba',
            'Predicciones'], loc='lower right')
plt.show()
```

C:\Users\Usuario\AppData\Local\Temp\ipykernel_14880\1792779524.py:4:
 SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 valid['Predictions'] = predictions



```
# Separar los datos de prueba
valid = data[training_data_len:]
valid['Predictions'] = predictions

# Graficar los datos de prueba y las predicciones
plt.figure(figsize=(16, 6))
plt.title('Google')
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre USD ($)', fontsize=18)
plt.plot(valid['Adj Close'], linewidth=1, label='Datos de prueba')
plt.plot(valid['Predictions'], linewidth=1, label='Predicciones')
plt.legend(loc='lower right')
plt.show()
```

C:\Users\Usuario\AppData\Local\Temp\ipykernel_14880\549889742.py:3:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
valid['Predictions'] = predictions



```
# Cargamos los datos de las acciones que queremos analizar y predecir
df = pdr.get_data_yahoo('MSFT', start='2012-01-01',
end=datetime.now())
```

df

[*****100%*****] 1 of 1 completed

\	Open	High	Low	Close	Adj Close
Date					
2012-01-03	26.549999	26.959999	26.389999	26.770000	21.200520
2012-01-04	26.820000	27.469999	26.780001	27.400000	21.699444
2012-01-05	27.379999	27.730000	27.290001	27.680000	21.921188
2012-01-06	27.530001	28.190001	27.530001	28.110001	22.261730
2012-01-09	28.049999	28.100000	27.719999	27.740000	21.968702
...
2024-05-13	418.010010	418.350006	410.820007	413.720001	412.975098
2024-05-14	412.019989	417.489990	411.549988	416.559998	415.809998
2024-05-15	417.899994	423.809998	417.269989	423.079987	423.079987
2024-05-16	421.799988	425.420013	420.350006	420.989990	420.989990
2024-05-17	422.170013	422.920013	418.024994	420.209991	420.209991

Volume

```

Date
2012-01-03    64731500
2012-01-04    80516100
2012-01-05    56081400
2012-01-06    99455500
2012-01-09    59706800
...
2024-05-13    15440200
2024-05-14    15109300
2024-05-15    22239500
2024-05-16    17530100
2024-05-17    14250616


```

```
[3114 rows x 6 columns]
```

```

#Hacemos una gráfica con el histórico de precios
plt.figure(figsize=(16,6))
plt.title('Histórico de precios MSFT')
plt.plot(df['Adj Close'], linewidth = 1)
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre ajustado USD ($)', fontsize=18)
plt.show()

```



```

# Creamos un nuevo DataFrame con únicamente la columna Adj Close
data = df.filter(['Adj Close'])
# Convertimos el dataframe en un numpy array
dataset = data.values
# Número de filas para entrenar el modelo
training_data_len = int(np.ceil( len(dataset) * .95 ))

training_data_len

```


2959

```
# Escalamos y normalizamos los datos
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler(feature_range=(0,1))
scaled_data = scaler.fit_transform(dataset)

scaled_data

array([[0.          ],
       [0.00122467],
       [0.00176896],
       ...,
       [0.98645806],
       [0.98132793],
       [0.97941334]])

# Creamos el data set para entrenamiento con los datos escalados
train_data = scaled_data[0:int(training_data_len), :]
# Separamos los datos entre x_train y y_train
x_train = []
y_train = []

for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:i, 0])
    y_train.append(train_data[i, 0])
    if i<= 61:
        print(x_train)
        print(y_train)
        print()

# Convertimos x_train y y_train en numpy arrays
x_train, y_train = np.array(x_train), np.array(y_train)

# Cambiamos la forma de los datos
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))

[array([0.          , 0.00122467, 0.00176896, 0.00260486, 0.00188559,
        0.00207998, 0.00184672, 0.002391  , 0.002877  , 0.00289645,
        0.00283814, 0.00262427, 0.00571513, 0.00575403, 0.00499588,
        0.00542357, 0.00530691, 0.00478206, 0.00552075, 0.00536523,
        0.00606504, 0.00618169, 0.00674543, 0.00666769, 0.00695926,
        0.00756189, 0.0077757  , 0.00725084, 0.00740636, 0.00715199,
        0.00676064, 0.00918698, 0.00910872, 0.00948049, 0.00914783,
        0.00934351, 0.00955876, 0.00930439, 0.01032189, 0.0100675 ,
        0.0111437  , 0.0107328  , 0.01018491, 0.00971529, 0.01026318,
        0.01059584, 0.0105567  , 0.01065454, 0.01188726, 0.01208295,
        0.01223948, 0.0117503  , 0.01096761, 0.0105567  , 0.01040015,
        0.01057625, 0.01059584, 0.01173073, 0.01159376, 0.01094803])]
```

```
[0.010811063577416284]
[array([0.          , 0.00122467, 0.00176896, 0.00260486, 0.00188559,
        0.00207998, 0.00184672, 0.002391  , 0.002877  , 0.00289645,
        0.00283814, 0.00262427, 0.00571513, 0.00575403, 0.00499588,
        0.00542357, 0.00530691, 0.00478206, 0.00552075, 0.00536523,
        0.00606504, 0.00618169, 0.00674543, 0.00666769, 0.00695926,
        0.00756189, 0.0077757  , 0.00725084, 0.00740636, 0.00715199,
        0.00676064, 0.00918698, 0.00910872, 0.00948049, 0.00914783,
        0.00934351, 0.00955876, 0.00930439, 0.01032189, 0.0100675  ,
        0.0111437  , 0.0107328  , 0.01018491, 0.00971529, 0.01026318,
        0.01059584, 0.0105567  , 0.01065454, 0.01188726, 0.01208295,
        0.01223948, 0.0117503  , 0.01096761, 0.0105567  , 0.01040015,
        0.01057625, 0.01059584, 0.01173073, 0.01159376, 0.01094803]),
array([0.00122467, 0.00176896, 0.00260486, 0.00188559, 0.00207998,
        0.00184672, 0.002391  , 0.002877  , 0.00289645, 0.00283814,
        0.00262427, 0.00571513, 0.00575403, 0.00499588, 0.00542357,
        0.00530691, 0.00478206, 0.00552075, 0.00536523, 0.00606504,
        0.00618169, 0.00674543, 0.00666769, 0.00695926, 0.00756189,
        0.0077757  , 0.00725084, 0.00740636, 0.00715199, 0.00676064,
        0.00918698, 0.00910872, 0.00948049, 0.00914783, 0.00934351,
        0.00955876, 0.00930439, 0.01032189, 0.0100675  , 0.0111437  ,
        0.0107328  , 0.01018491, 0.00971529, 0.01026318, 0.01059584,
        0.0105567  , 0.01065454, 0.01188726, 0.01208295, 0.01223948,
        0.0117503  , 0.01096761, 0.0105567  , 0.01040015, 0.01057625,
        0.01059584, 0.01173073, 0.01159376, 0.01094803, 0.01081106])],
[0.010811063577416284, 0.011085009761975122])
```

```
from keras.models import Sequential
from keras.layers import LSTM, Dense

# Creamos el modelo LSTM
model = Sequential()

# Añadimos capas LSTM con 50 unidades
model.add(LSTM(50, return_sequences=True,
input_shape=(x_train.shape[1], 1)))
model.add(LSTM(50, return_sequences=False))
# Añadimos una capa densa con 25 unidades
model.add(Dense(25))
# Añadir una capa de salida densa con 1 unidad
model.add(Dense(1))

# Compilamos el modelo
model.compile(optimizer='adam', loss='mean_squared_error')

# Entrenamos el modelo
model.fit(x_train, y_train, batch_size=1, epochs=1)
```

```

# Creamos el conjunto de datos de prueba
test_data = scaled_data[training_data_len - 60: , :]
# Creamos los conjuntos de datos x_test y y_test
x_test = []
y_test = dataset[training_data_len:, :]
for i in range(60, len(test_data)):
    x_test.append(test_data[i-60:i, 0])

# Convertimos a un numpy array
x_test = np.array(x_test)

# Cambiamos la forma de los datos
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))

# Obtención de las predicciones
predictions = model.predict(x_test)
predictions = scaler.inverse_transform(predictions)

# Calculamos el RMSE
rmse=np.sqrt(np.mean(((predictions- y_test)**2)))
print("Root Mean Square Error:",rmse)

# Calcular el NRMSE
range_y = np.max(y_test) - np.min(y_test)
nrmse = rmse / range_y
print("Normalized Root Mean Square Error (NRMSE):", nrmse)

# Calculamos el MAE
mae = np.mean(np.abs(predictions - y_test))
print("Mean Absolute Error:", mae)

# Calculamos el coeficiente de determinación (R cuadrado)
from sklearn.metrics import r2_score

r2 = r2_score(y_test, predictions)
print("Coeficiente de determinación (R cuadrado):", r2)

```

```

c:\Users\Usuario\anaconda3\Lib\site-packages\keras\src\layers\rnn\
rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
    super().__init__(**kwargs)

```

```

2899/2899 ————— 25s 8ms/step - loss: 0.0012
5/5 ————— 0s 57ms/step
Root Mean Square Error: 10.859410012160629
Normalized Root Mean Square Error (NRMSE): 0.1046425412665281
Mean Absolute Error: 9.063875063004032
Coeficiente de determinación (R cuadrado): 0.8621789920178538

```

```
# Graficamos los datos
```

```
train = data[:training_data_len]  
valid = data[training_data_len:]  
valid['Predictions'] = predictions
```

```
plt.figure(figsize=(16,6))  
plt.title('Microsoft')  
plt.xlabel('Fecha', fontsize=18)  
plt.ylabel('Cierre USD ($)', fontsize=18)  
plt.plot(train['Adj Close'], linewidth = 1)  
plt.plot(valid[['Adj Close', 'Predictions']], linewidth = 1)  
plt.legend(['Datos de entrenamiento', 'Datos de prueba',  
            'Predicciones'], loc='lower right')  
plt.show()
```


C:\Users\Usuario\AppData\Local\Temp\ipykernel_14880\3267775332.py:4:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
valid['Predictions'] = predictions
```



```
# Separar los datos de prueba
```

```
valid = data[training_data_len:]  
valid['Predictions'] = predictions
```


```
# Graficar los datos de prueba y las predicciones
```

```
plt.figure(figsize=(16, 6))  
plt.title('Microsoft')  
plt.xlabel('Fecha', fontsize=18)  
plt.ylabel('Cierre USD ($)', fontsize=18)
```

```
plt.plot(valid['Adj Close'], linewidth=1, label='Datos de prueba')
plt.plot(valid['Predictions'], linewidth=1, label='Predicciones')
plt.legend(loc='lower right')
plt.show()
```

C:\Users\Usuario\AppData\Local\Temp\ipykernel_14880\1228714177.py:3:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
valid['Predictions'] = predictions



```
# Cargamos los datos de las acciones que queremos analizar y predecir
df = pdr.get_data_yahoo('META', start='2012-01-01',
end=datetime.now())
```

df

[*****100%*****] 1 of 1 completed

\	Open	High	Low	Close	Adj Close
Date					
2012-05-18	42.049999	45.000000	38.000000	38.230000	38.189480
2012-05-21	36.529999	36.660000	33.000000	34.029999	33.993931
2012-05-22	32.610001	33.590000	30.940001	31.000000	30.967144
2012-05-23	31.370001	32.500000	31.360001	32.000000	31.966084

2012-05-24	32.950001	33.209999	31.770000	33.029999	32.994991
...
2024-05-13	472.750000	473.350006	462.850006	468.010010	468.010010
2024-05-14	463.369995	472.540009	460.079987	471.850006	471.850006
2024-05-15	474.980011	482.500000	471.200012	481.540009	481.540009
2024-05-16	475.000000	477.690002	472.750000	473.230011	473.230011
2024-05-17	470.589996	472.250000	468.420013	471.910004	471.910004

	Volume
Date	
2012-05-18	573576400
2012-05-21	168192700
2012-05-22	101786600
2012-05-23	73600000
2012-05-24	50237200
...	...
2024-05-13	14668800
2024-05-14	10478600
2024-05-15	13100500
2024-05-16	16608200
2024-05-17	9964026

[3019 rows x 6 columns]

```
#Hacemos una gráfica con el histórico de precios
plt.figure(figsize=(16,6))
plt.title('Histórico de precios META')
plt.plot(df['Adj Close'], linewidth = 1)
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre ajustado USD ($)', fontsize=18)
plt.show()
```



```
# Creamos un nuevo DataFrame con únicamente la columna Adj Close
data = df.filter(['Adj Close'])
# Convertimos el dataframe en un numpy array
dataset = data.values
# Número de filas para entrenar el modelo
training_data_len = int(np.ceil( len(dataset) * .95 ))
```

```
training_data_len
```

```
2869
```

```
# Escalamos y normalizamos los datos
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler(feature_range=(0,1))
scaled_data = scaler.fit_transform(dataset)
```

```
scaled_data
```

```
array([[0.04018272],
       [0.03195016],
       [0.02601096],
       ...,
       [0.91013063],
       [0.89382465],
       [0.89123452]])
```

```
# Creamos el data set para entrenamiento con los datos escalados
train_data = scaled_data[0:int(training_data_len), :]
# Separamos los datos entre x_train y y_train
x_train = []
y_train = []
```

```

for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:i, 0])
    y_train.append(train_data[i, 0])
    if i<= 61:
        print(x_train)
        print(y_train)
        print()

# Convertimos x_train y y_train en numpy arrays
x_train, y_train = np.array(x_train), np.array(y_train)

# Cambiamos la forma de los datos
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))

[array([0.04018272, 0.03195016, 0.02601096, 0.02797109, 0.02999003,
        0.02779468, 0.02177707, 0.02050299, 0.02326678, 0.01958172,
        0.01797442, 0.01595548, 0.017798 , 0.01681794, 0.01836644,
        0.01819003, 0.01895448, 0.01869966, 0.020699 , 0.02407043,
        0.02681462, 0.02779468, 0.02718704, 0.02765747, 0.03002923,
        0.0280887 , 0.03012723, 0.02842192, 0.02671661, 0.02620697,
        0.02556013, 0.02640299, 0.02693222, 0.02744186, 0.02830431,
        0.02693222, 0.02595215, 0.02563853, 0.02546212, 0.0206206 ,
        0.02030698, 0.02230631, 0.02209069, 0.02162026, 0.02160066,
        0.02101262, 0.02275714, 0.01787641, 0.01172159, 0.01062392,
        0.00780133, 0.00617442, 0.00452791, 0.00658605, 0.00821295,
        0.00586079, 0.00586079, 0.00642924, 0.00799734, 0.00758571])]
[0.0051943492946955536]

[array([0.04018272, 0.03195016, 0.02601096, 0.02797109, 0.02999003,
        0.02779468, 0.02177707, 0.02050299, 0.02326678, 0.01958172,
        0.01797442, 0.01595548, 0.017798 , 0.01681794, 0.01836644,
        0.01819003, 0.01895448, 0.01869966, 0.020699 , 0.02407043,
        0.02681462, 0.02779468, 0.02718704, 0.02765747, 0.03002923,
        0.0280887 , 0.03012723, 0.02842192, 0.02671661, 0.02620697,
        0.02556013, 0.02640299, 0.02693222, 0.02744186, 0.02830431,
        0.02693222, 0.02595215, 0.02563853, 0.02546212, 0.0206206 ,
        0.02030698, 0.02230631, 0.02209069, 0.02162026, 0.02160066,
        0.02101262, 0.02275714, 0.01787641, 0.01172159, 0.01062392,
        0.00780133, 0.00617442, 0.00452791, 0.00658605, 0.00821295,
        0.00586079, 0.00586079, 0.00642924, 0.00799734, 0.00758571]),
array([0.03195016, 0.02601096, 0.02797109, 0.02999003, 0.02779468,
        0.02177707, 0.02050299, 0.02326678, 0.01958172, 0.01797442,
        0.01595548, 0.017798 , 0.01681794, 0.01836644, 0.01819003,
        0.01895448, 0.01869966, 0.020699 , 0.02407043, 0.02681462,
        0.02779468, 0.02718704, 0.02765747, 0.03002923, 0.0280887 ,
        0.03012723, 0.02842192, 0.02671661, 0.02620697, 0.02556013,
        0.02640299, 0.02693222, 0.02744186, 0.02830431, 0.02693222,
        0.02595215, 0.02563853, 0.02546212, 0.0206206 , 0.02030698,
        0.02230631, 0.02209069, 0.02162026, 0.02160066, 0.02101262,
        0.02275714, 0.01787641, 0.01172159, 0.01062392, 0.00780133,

```



```
0.00617442, 0.00452791, 0.00658605, 0.00821295, 0.00586079,  
0.00586079, 0.00642924, 0.00799734, 0.00758571, 0.00519435]])  
[0.0051943492946955536, 0.006801659991111868]
```

```
from keras.models import Sequential  
from keras.layers import LSTM, Dense  
  
# Creamos el modelo LSTM  
model = Sequential()  
  
# Añadimos capas LSTM con 50 unidades  
model.add(LSTM(50, return_sequences=True,  
input_shape=(x_train.shape[1], 1)))  
model.add(LSTM(50, return_sequences=False))  
# Añadimos una capa densa con 25 unidades  
model.add(Dense(25))  
# Añadir una capa de salida densa con 1 unidad  
model.add(Dense(1))  
  
# Compilamos el modelo  
model.compile(optimizer='adam', loss='mean_squared_error')  
  
# Entrenamos el modelo  
model.fit(x_train, y_train, batch_size=1, epochs=1)  
  
# Creamos el conjunto de datos de prueba  
test_data = scaled_data[training_data_len - 60: , :]  
# Creamos los conjuntos de datos x_test y y_test  
x_test = []  
y_test = dataset[training_data_len:, :]  
for i in range(60, len(test_data)):  
    x_test.append(test_data[i-60:i, 0])  
  
# Convertimos a un numpy array  
x_test = np.array(x_test)  
  
# Cambiamos la forma de los datos  
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))  
  
# Obtención de las predicciones  
predictions = model.predict(x_test)  
predictions = scaler.inverse_transform(predictions)  
  
# Calculamos el RMSE  
rmse=np.sqrt(np.mean(((predictions- y_test)**2)))  
print("Root Mean Square Error:",rmse)  
  
# Calcular el NRMSE  
range_y = np.max(y_test) - np.min(y_test)  
nrmse = rmse / range_y
```

```

print("Normalized Root Mean Square Error (NRMSE):", nrmse)

# Calculamos el MAE
mae = np.mean(np.abs(predictions - y_test))
print("Mean Absolute Error:", mae)

# Calculamos el coeficiente de determinación (R cuadrado)
from sklearn.metrics import r2_score

r2 = r2_score(y_test, predictions)
print("Coeficiente de determinación (R cuadrado):", r2)

c:\Users\Usuario\anaconda3\Lib\site-packages\keras\src\layers\rnn\
rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

```

```

2809/2809 ————— 24s 8ms/step - loss: 0.0016
5/5 ————— 0s 57ms/step
Root Mean Square Error: 18.694869492613222
Normalized Root Mean Square Error (NRMSE): 0.07812456748345288
Mean Absolute Error: 14.901309611002604
Coeficiente de determinación (R cuadrado): 0.9391195888533523

```

```

# Graficamos los datos
train = data[:training_data_len]
valid = data[training_data_len:]
valid['Predictions'] = predictions

```

```

plt.figure(figsize=(16,6))
plt.title('META')
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre USD ($)', fontsize=18)
plt.plot(train['Adj Close'], linewidth = 1)
plt.plot(valid[['Adj Close', 'Predictions']], linewidth = 1)
plt.legend(['Datos de entrenamiento', 'Datos de prueba',
'Predicciones'], loc='lower right')
plt.show()

```

```


C:\Users\Usuario\AppData\Local\Temp\ipykernel_14880\474930955.py:4:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

```

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#
returning-a-view-versus-a-copy
valid['Predictions'] = predictions

```




```
# Separar los datos de prueba
valid = data[training_data_len:]
valid['Predictions'] = predictions

# Graficar los datos de prueba y las predicciones
plt.figure(figsize=(16, 6))
plt.title('META')
plt.xlabel('Fecha', fontsize=18)
plt.ylabel('Cierre USD ($)', fontsize=18)
plt.plot(valid['Adj Close'], linewidth=1, label='Datos de prueba')
plt.plot(valid['Predictions'], linewidth=1, label='Predicciones')
plt.legend(loc='lower right')
plt.show()
```

C:\Users\Usuario\AppData\Local\Temp\ipykernel_14880\2918790472.py:3:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
valid['Predictions'] = predictions



Predicción con XGBoost

```
df2 = pdr.get_data_yahoo('AAPL', start='2012-01-01',
end=datetime.now())
Datos = df2.filter(["Date", "Adj Close"])
Datos.head()

[*****100%*****] 1 of 1 completed
```

Date	Adj Close
2012-01-03	12.416985
2012-01-04	12.483712
2012-01-05	12.622309
2012-01-06	12.754258
2012-01-09	12.734030

```
import numpy as np
import xgboost as xgb
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error
import pandas as pd

# Calculamos la diferencia de precios como una nueva característica
# para conseguir una predicción más precisa
Datos['Price_Diff'] = Datos['Adj Close'].diff()

# Eliminamos la primera fila que contendría un valor NaN después del
# cálculo de la diferencia
datos = Datos.dropna()

# Separamos los datos en conjuntos de entrenamiento y prueba
training_data_len = int(len(datos) * 0.95)
training_data = datos.iloc[:training_data_len]
test_data = datos.iloc[training_data_len:]
```

```

# Convertimos las fechas en características numéricas
X_train = np.array(range(len(training_data)).reshape(-1, 1))
X_train = np.hstack((X_train,
training_data['Price_Diff'].values.reshape(-1, 1)))
y_train = training_data['Adj Close'].values

# Iniciamos el modelo XGBoost
model_xgb = xgb.XGBRegressor(objective='reg:squarederror')

# Definimos los parámetros a ajustar
param_grid = {
    'max_depth': [3, 4, 5, 6, 7],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'n_estimators': [100, 200, 300, 400]
}

# Iniciamos el GridSearchCV
grid_search = GridSearchCV(estimator=model_xgb, param_grid=param_grid,
                           scoring='neg_root_mean_squared_error',
                           cv=3, n_jobs=-1, verbose=2)

# Realizamos la búsqueda de los mejores hiperparámetros
grid_search.fit(X_train, y_train)

# Obtenemos los mejores hiperparámetros
best_params = grid_search.best_params_
print("Mejores hiperparámetros:", best_params)

# Obtenemos el mejor modelo
best_model = grid_search.best_estimator_

# Convertimos las fechas de prueba en características numéricas
X_test = np.array(range(training_data_len, len(datos)).reshape(-1, 1))
X_test = np.hstack((X_test, test_data['Price_Diff'].values.reshape(-1,
1)))
y_test = test_data['Adj Close'].values

# Hacemos predicciones en el conjunto de prueba
predictions = best_model.predict(X_test)

# Calculamos el RMSE
rmse = np.sqrt(mean_squared_error(y_test, predictions))
print("Root Mean Square Error (XGBoost):", rmse)

# Calcular el NRMSE
range_y = np.max(y_test) - np.min(y_test)
nrmse = rmse / range_y
print("Normalized Root Mean Square Error (XGBoost):", nrmse)

```

```

# Calculamos el MAE
mae = np.mean(np.abs(predictions - y_test))
print("Mean Absolute Error:", mae)

# Calculamos el coeficiente de determinación (R cuadrado)
from sklearn.metrics import r2_score

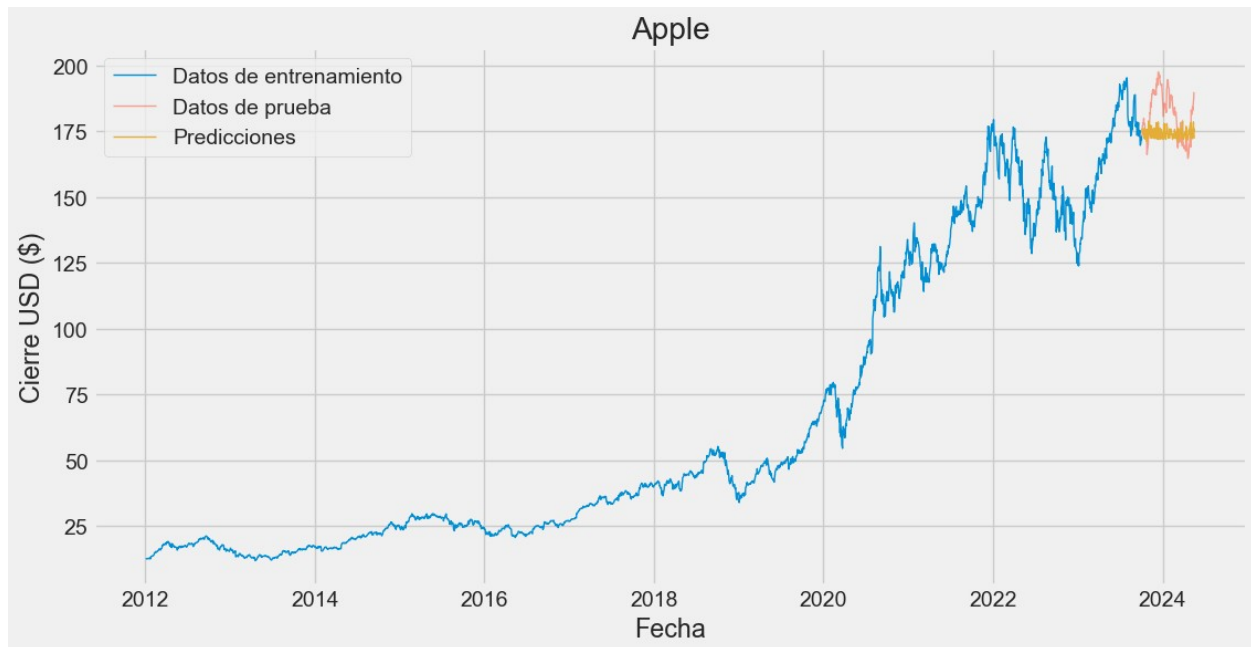
r2 = r2_score(y_test, predictions)
print("Coeficiente de determinación (R cuadrado):", r2)

Fitting 3 folds for each of 100 candidates, totalling 300 fits
Mejores hiperparámetros: {'learning_rate': 0.1, 'max_depth': 4,
'n_estimators': 200}
Root Mean Square Error (XGBoost): 11.093225389453506
Normalized Root Mean Square Error (XGBoost): 0.33807390467376586
Mean Absolute Error: 9.233179141313602
Coeficiente de determinación (R cuadrado): -0.5454802841811077


import matplotlib.pyplot as plt

# Graficamos los datos
plt.figure(figsize=(12, 6))
plt.plot(training_data.index, training_data['Adj Close'], label='Datos
de entrenamiento', linewidth = 1)
plt.plot(test_data.index, test_data['Adj Close'], label='Datos de
prueba', alpha=0.5, linewidth = 1)
plt.plot(test_data.index, predictions, label='Predicciones',
linestyle='--', linewidth = 1)
plt.title('Apple')
plt.xlabel('Fecha')
plt.ylabel('Cierre USD ($)')
plt.legend()
plt.grid(True)
plt.show()

```



```
plt.figure(figsize=(12, 6))
plt.plot(test_data.index, test_data['Adj Close'], label='Datos de prueba', alpha=0.5, linewidth=1)
plt.plot(test_data.index, predictions, label='Predicciones', linestyle='--', linewidth=1)
plt.title('Apple')
plt.xlabel('Fecha')
plt.ylabel('Cierre USD ($)')
plt.legend()
plt.grid(True)
plt.show()
```



```
df2 = pdr.get_data_yahoo('GOOG', start='2012-01-01',
end=datetime.now())
Datos = df2.filter(["Date", "Adj Close"])
Datos.head()
```

```
[*****100%*****] 1 of 1 completed
```

Date	Adj Close
2012-01-03	16.573130
2012-01-04	16.644611
2012-01-05	16.413727
2012-01-06	16.189817
2012-01-09	15.503389

```
import numpy as np
import xgboost as xgb
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error
import pandas as pd
```

```
# Calculamos la diferencia de precios como una nueva característica
para conseguir una predicción más precisa
```

```
Datos['Price_Diff'] = Datos['Adj Close'].diff()
```

```
# Eliminamos la primera fila que contendría un valor NaN después del
cálculo de la diferencia
```

```
datos = Datos.dropna()
```

```
# Separamos los datos en conjuntos de entrenamiento y prueba
```



```

training_data_len = int(len(datos) * 0.95)
training_data = datos.iloc[:training_data_len]
test_data = datos.iloc[training_data_len:]

# Convertimos las fechas en características numéricas
X_train = np.array(range(len(training_data))).reshape(-1, 1)
X_train = np.hstack((X_train,
training_data['Price_Diff'].values.reshape(-1, 1)))
y_train = training_data['Adj Close'].values

# Iniciamos el modelo XGBoost
model_xgb = xgb.XGBRegressor(objective='reg:squarederror')

# Definimos los parámetros a ajustar
param_grid = {
    'max_depth': [3, 4, 5, 6, 7],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'n_estimators': [100, 200, 300]
}

# Iniciamos el GridSearchCV
grid_search = GridSearchCV(estimator=model_xgb, param_grid=param_grid,
                           scoring='neg_root_mean_squared_error',
                           cv=3, n_jobs=-1, verbose=2)

# Realizamos la búsqueda de los mejores hiperparámetros
grid_search.fit(X_train, y_train)

# Obtenemos los mejores hiperparámetros
best_params = grid_search.best_params_
print("Mejores hiperparámetros:", best_params)

# Obtenemos el mejor modelo
best_model = grid_search.best_estimator_

# Convertimos las fechas de prueba en características numéricas
X_test = np.array(range(training_data_len, len(datos))).reshape(-1, 1)
X_test = np.hstack((X_test, test_data['Price_Diff'].values.reshape(-1,
1)))
y_test = test_data['Adj Close'].values

# Hacemos predicciones en el conjunto de prueba
predictions = best_model.predict(X_test)

# Calculamos el RMSE
rmse = np.sqrt(mean_squared_error(y_test, predictions))
print("Root Mean Square Error (XGBoost):", rmse)

# Calcular el NRMSE
range_y = np.max(y_test) - np.min(y_test)

```

```
nrmse = rmse / range_y  
print("Normalized Root Mean Square Error (XGBoost):", nrmse)
```

```
# Calculamos el MAE
```

```
mae = np.mean(np.abs(predictions - y_test))  
print("Mean Absolute Error:", mae)
```

```
# Calculamos el coeficiente de determinación (R cuadrado)
```

```
from sklearn.metrics import r2_score
```

```
r2 = r2_score(y_test, predictions)
```

```
print("Coeficiente de determinación (R cuadrado):", r2)
```

Fitting 3 folds for each of 75 candidates, totalling 225 fits

Mejores hiperparámetros: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 300}

Root Mean Square Error (XGBoost): 15.776395795843193

Normalized Root Mean Square Error (XGBoost): 0.3003883626090438

Mean Absolute Error: 11.83697754297501

Coeficiente de determinación (R cuadrado): -0.8287643671716405

```
import matplotlib.pyplot as plt
```

```
# Graficamos los datos
```

```
plt.figure(figsize=(12, 6))
```

```
plt.plot(training_data.index, training_data['Adj Close'], label='Datos  
de entrenamiento', linewidth = 1)
```

```
plt.plot(test_data.index, test_data['Adj Close'], label='Datos de  
prueba', alpha=0.5, linewidth = 1)
```

```
plt.plot(test_data.index, predictions, label='Predicciones',  
linestyle='--', linewidth = 1)
```

```
plt.title('Google')
```


```
plt.xlabel('Fecha')
```

```
plt.ylabel('Cierre USD ($)')
```


```
plt.legend()
```

```
plt.grid(True)
```

```
plt.show()
```



```
plt.figure(figsize=(12, 6))
plt.plot(test_data.index, test_data['Adj Close'], label='Datos de prueba', alpha=0.5, linewidth=1)
plt.plot(test_data.index, predictions, label='Predicciones', linestyle='--', linewidth=1)
plt.title('Google')
plt.xlabel('Fecha')
plt.ylabel('Cierre USD ($)')
plt.legend()
plt.grid(True)
plt.show()
```



```
df2 = pdr.get_data_yahoo('MSFT', start='2012-01-01',
end=datetime.now())
Datos = df2.filter(["Date", "Adj Close"])
Datos.head()
```

```
[*****100%*****] 1 of 1 completed
```

Date	Adj Close
2012-01-03	21.200516
2012-01-04	21.699440
2012-01-05	21.921190
2012-01-06	22.261724
2012-01-09	21.968702

```
import numpy as np
import xgboost as xgb
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error
import pandas as pd
```

```
# Calculamos la diferencia de precios como una nueva característica
para conseguir una predicción más precisa
```

```
Datos['Price_Diff'] = Datos['Adj Close'].diff()
```

```
# Eliminamos la primera fila que contendría un valor NaN después del
cálculo de la diferencia
```

```
datos = Datos.dropna()
```

```
# Separamos los datos en conjuntos de entrenamiento y prueba
```

```
training_data_len = int(len(datos) * 0.95)
```

```

training_data = datos.iloc[:training_data_len]
test_data = datos.iloc[training_data_len:]

# Convertimos las fechas en características numéricas
X_train = np.array(range(len(training_data))).reshape(-1, 1)
X_train = np.hstack((X_train,
training_data['Price_Diff'].values.reshape(-1, 1)))
y_train = training_data['Adj Close'].values

# Iniciamos el modelo XGBoost
model_xgb = xgb.XGBRegressor(objective='reg:squarederror')

# Definimos los parámetros a ajustar
param_grid = {
    'max_depth': [3, 4, 5, 6, 7],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'n_estimators': [100, 200, 300]
}

# Iniciamos el GridSearchCV
grid_search = GridSearchCV(estimator=model_xgb, param_grid=param_grid,
                           scoring='neg_root_mean_squared_error',
                           cv=3, n_jobs=-1, verbose=2)

# Realizamos la búsqueda de los mejores hiperparámetros
grid_search.fit(X_train, y_train)

# Obtenemos los mejores hiperparámetros
best_params = grid_search.best_params_
print("Mejores hiperparámetros:", best_params)

# Obtenemos el mejor modelo
best_model = grid_search.best_estimator_

# Convertimos las fechas de prueba en características numéricas
X_test = np.array(range(training_data_len, len(datos))).reshape(-1, 1)
X_test = np.hstack((X_test, test_data['Price_Diff'].values.reshape(-1,
1)))
y_test = test_data['Adj Close'].values

# Hacemos predicciones en el conjunto de prueba
predictions = best_model.predict(X_test)

# Calculamos el RMSE
rmse = np.sqrt(mean_squared_error(y_test, predictions))
print("Root Mean Square Error (XGBoost):", rmse)

# Calcular el NRMSE
range_y = np.max(y_test) - np.min(y_test)
nrmse = rmse / range_y

```

```

print("Normalized Root Mean Square Error (XGBoost):", nrmse)

# Calculamos el MAE
mae = np.mean(np.abs(predictions - y_test))
print("Mean Absolute Error:", mae)

# Calculamos el coeficiente de determinación (R cuadrado)
from sklearn.metrics import r2_score

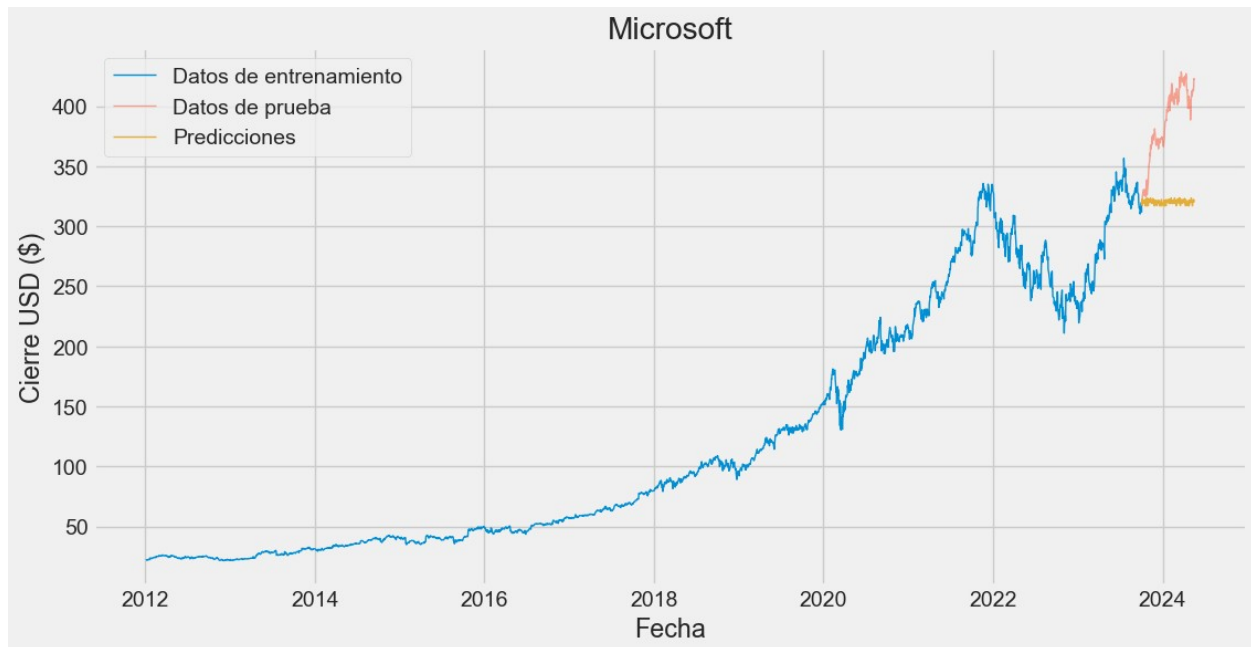
r2 = r2_score(y_test, predictions)
print("Coeficiente de determinación (R cuadrado):", r2)

Fitting 3 folds for each of 75 candidates, totalling 225 fits
Mejores hiperparámetros: {'learning_rate': 0.05, 'max_depth': 3,
'n_estimators': 200}
Root Mean Square Error (XGBoost): 72.83048199720946
Normalized Root Mean Square Error (XGBoost): 0.6535246335482785
Mean Absolute Error: 66.44710090832832
Coeficiente de determinación (R cuadrado): -4.853648411061348


import matplotlib.pyplot as plt

# Graficamos los datos
plt.figure(figsize=(12, 6))
plt.plot(training_data.index, training_data['Adj Close'], label='Datos
de entrenamiento', linewidth = 1)
plt.plot(test_data.index, test_data['Adj Close'], label='Datos de
prueba', alpha=0.5, linewidth = 1)
plt.plot(test_data.index, predictions, label='Predicciones',
linestyle='--', linewidth = 1)
plt.title('Microsoft')
plt.xlabel('Fecha')
plt.ylabel('Cierre USD ($)')
plt.legend()
plt.grid(True)
plt.show()

```



```
plt.figure(figsize=(12, 6))
plt.plot(test_data.index, test_data['Adj Close'], label='Datos de prueba', alpha=0.5, linewidth=1)
plt.plot(test_data.index, predictions, label='Predicciones', linestyle='--', linewidth=1)
plt.title('Microsoft')
plt.xlabel('Fecha')
plt.ylabel('Cierre USD ($)')
plt.legend()
plt.grid(True)
plt.show()
```



```
df2 = pdr.get_data_yahoo('META', start='2012-01-01',
end=datetime.now())
Datos = df2.filter(["Date", "Adj Close"])
Datos.head()
```

[*****100%*****] 1 of 1 completed

Date	Adj Close
2012-05-18	38.189480
2012-05-21	33.993931
2012-05-22	30.967144
2012-05-23	31.966084
2012-05-24	32.994991

```
import numpy as np
import xgboost as xgb
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error
import pandas as pd
```

Calculamos la diferencia de precios como una nueva característica para conseguir una predicción más precisa

```
Datos['Price_Diff'] = Datos['Adj Close'].diff()
```

Eliminamos la primera fila que contendría un valor NaN después del cálculo de la diferencia

```
datos = Datos.dropna()
```

Separamos los datos en conjuntos de entrenamiento y prueba


```

training_data_len = int(len(datos) * 0.95)
training_data = datos.iloc[:training_data_len]
test_data = datos.iloc[training_data_len:]

# Convertimos las fechas en características numéricas
X_train = np.array(range(len(training_data))).reshape(-1, 1)
X_train = np.hstack((X_train,
training_data['Price_Diff'].values.reshape(-1, 1)))
y_train = training_data['Adj Close'].values

# Iniciamos el modelo XGBoost
model_xgb = xgb.XGBRegressor(objective='reg:squarederror')

# Definimos los parámetros a ajustar
param_grid = {
    'max_depth': [3, 4, 5, 6, 7],
    'learning_rate': [0.01, 0.05, 0.1, 0.2, 1],
    'n_estimators': [100, 200, 300, 1100]
}

# Iniciamos el GridSearchCV
grid_search = GridSearchCV(estimator=model_xgb, param_grid=param_grid,
                           scoring='neg_root_mean_squared_error',
                           cv=3, n_jobs=-1, verbose=2)

# Realizamos la búsqueda de los mejores hiperparámetros
grid_search.fit(X_train, y_train)

# Obtenemos los mejores hiperparámetros
best_params = grid_search.best_params_
print("Mejores hiperparámetros:", best_params)

# Obtenemos el mejor modelo
best_model = grid_search.best_estimator_

# Convertimos las fechas de prueba en características numéricas
X_test = np.array(range(training_data_len, len(datos))).reshape(-1, 1)
X_test = np.hstack((X_test, test_data['Price_Diff'].values.reshape(-1,
1)))
y_test = test_data['Adj Close'].values

# Hacemos predicciones en el conjunto de prueba
predictions = best_model.predict(X_test)

# Calculamos el RMSE
rmse = np.sqrt(mean_squared_error(y_test, predictions))
print("Root Mean Square Error (XGBoost):", rmse)

# Calcular el NRMSE
range_y = np.max(y_test) - np.min(y_test)

```

```
nrmse = rmse / range_y
print("Normalized Root Mean Square Error (XGBoost):", nrmse)
```

```
# Calculamos el MAE
```

```
mae = np.mean(np.abs(predictions - y_test))
print("Mean Absolute Error:", mae)
```

```
# Calculamos el coeficiente de determinación (R cuadrado)
```

```
from sklearn.metrics import r2_score
```

```
r2 = r2_score(y_test, predictions)
```

```
print("Coeficiente de determinación (R cuadrado):", r2)
```

Fitting 3 folds for each of 120 candidates, totalling 360 fits
Mejores hiperparámetros: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 1100}

Root Mean Square Error (XGBoost): 130.3500841904394

Normalized Root Mean Square Error (XGBoost): 0.5447239924747002

Mean Absolute Error: 106.53563200085368

Coeficiente de determinación (R cuadrado): -1.9412845051687757

```
import matplotlib.pyplot as plt
```

```
# Graficamos los datos
```

```
plt.figure(figsize=(12, 6))
```

```
plt.plot(training_data.index, training_data['Adj Close'], label='Datos de entrenamiento', linewidth = 1)
```

```
plt.plot(test_data.index, test_data['Adj Close'], label='Datos de prueba', alpha=0.5, linewidth = 1)
```

```
plt.plot(test_data.index, predictions, label='Predicciones', linestyle='-', linewidth = 1)
```

```
plt.title('META')
```

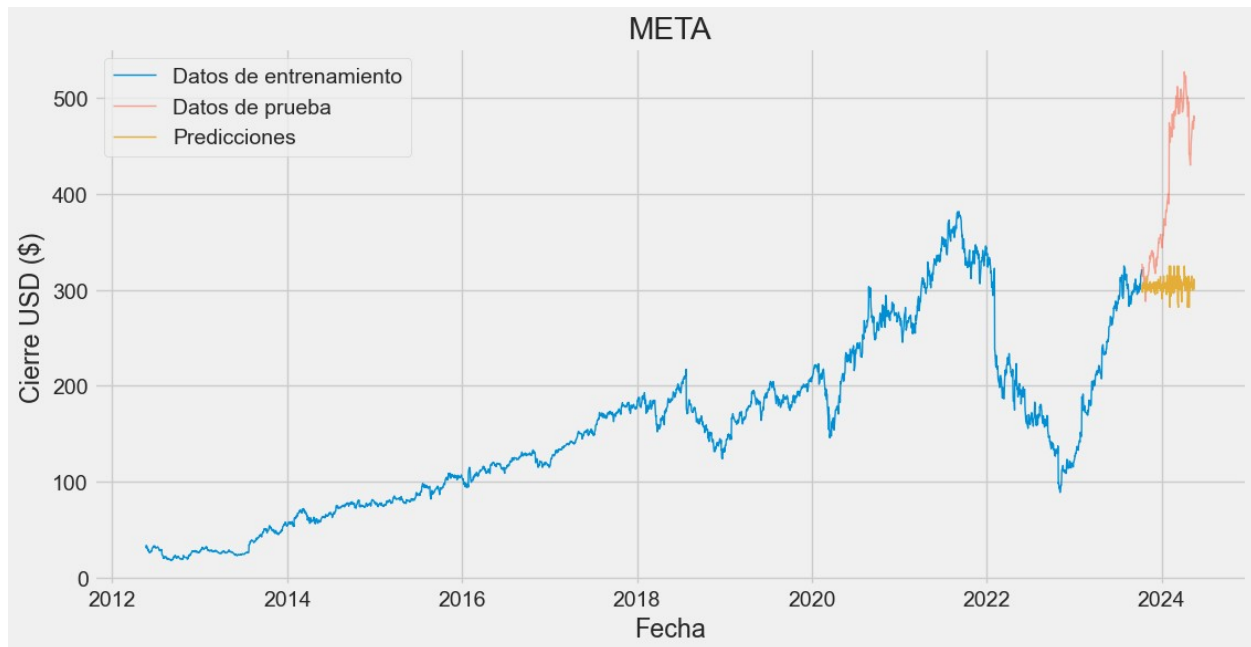
```
plt.xlabel('Fecha')
```

```
plt.ylabel('Cierre USD ($)')
```

```
plt.legend()
```

```
plt.grid(True)
```

```
plt.show()
```



```
plt.figure(figsize=(12, 6))
plt.plot(test_data.index, test_data['Adj Close'], label='Datos de prueba', alpha=0.5, linewidth=1)
plt.plot(test_data.index, predictions, label='Predicciones', linestyle='--', linewidth=1)
plt.title('META')
plt.xlabel('Fecha')
plt.ylabel('Cierre USD ($)')
plt.legend()
plt.grid(True)
plt.show()
```

